

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
import warnings
warnings.filterwarnings('ignore')

df= pd.read_csv("Mall_Customers.csv")
df.head()
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

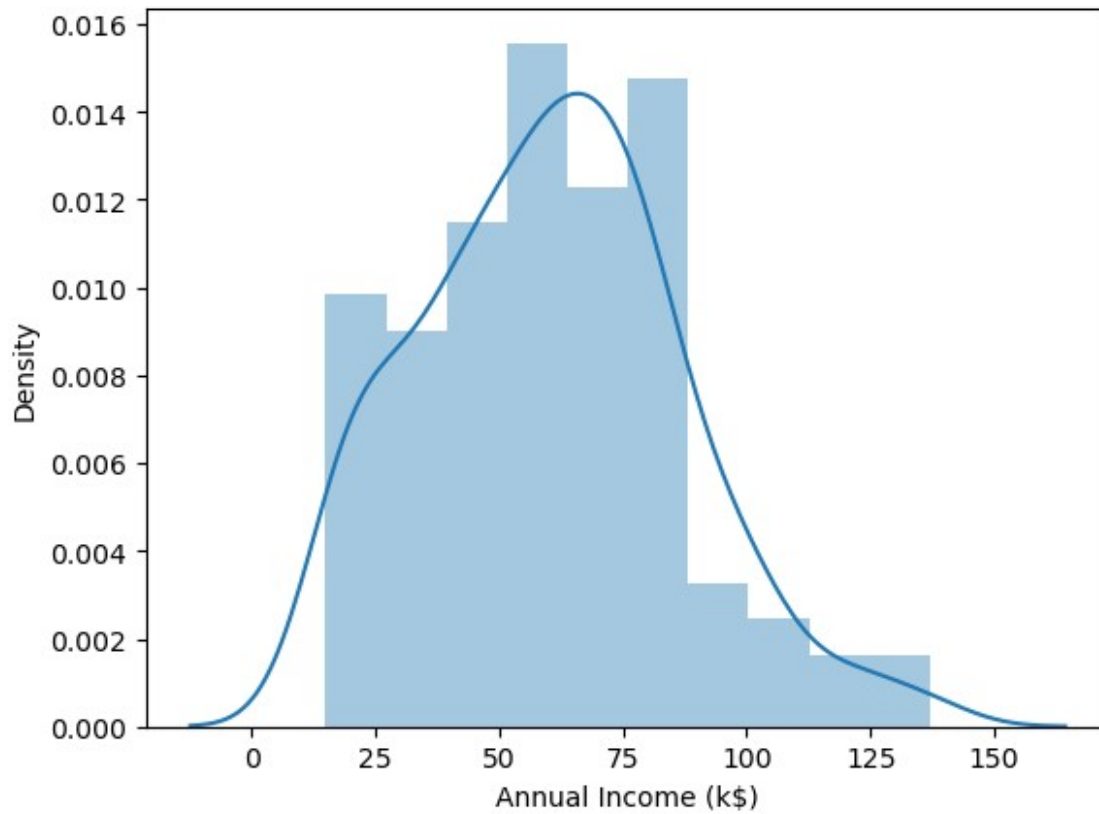
## Univariate Analysis

```
df.describe()
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
count	200.000000	200.000000	200.000000	200.000000
mean	100.500000	38.850000	60.560000	50.200000
std	57.879185	13.969007	26.264721	25.823522
min	1.000000	18.000000	15.000000	1.000000
25%	50.750000	28.750000	41.500000	34.750000
50%	100.500000	36.000000	61.500000	50.000000
75%	150.250000	49.000000	78.000000	73.000000
max	200.000000	70.000000	137.000000	99.000000

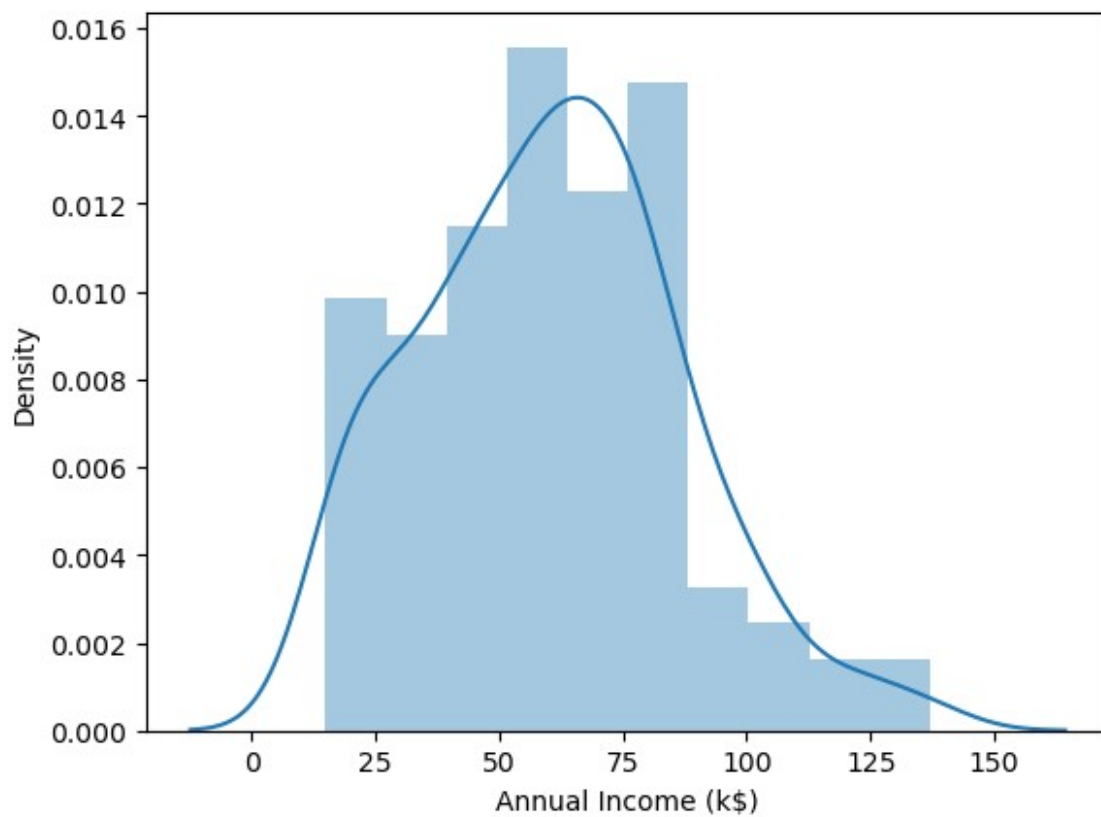
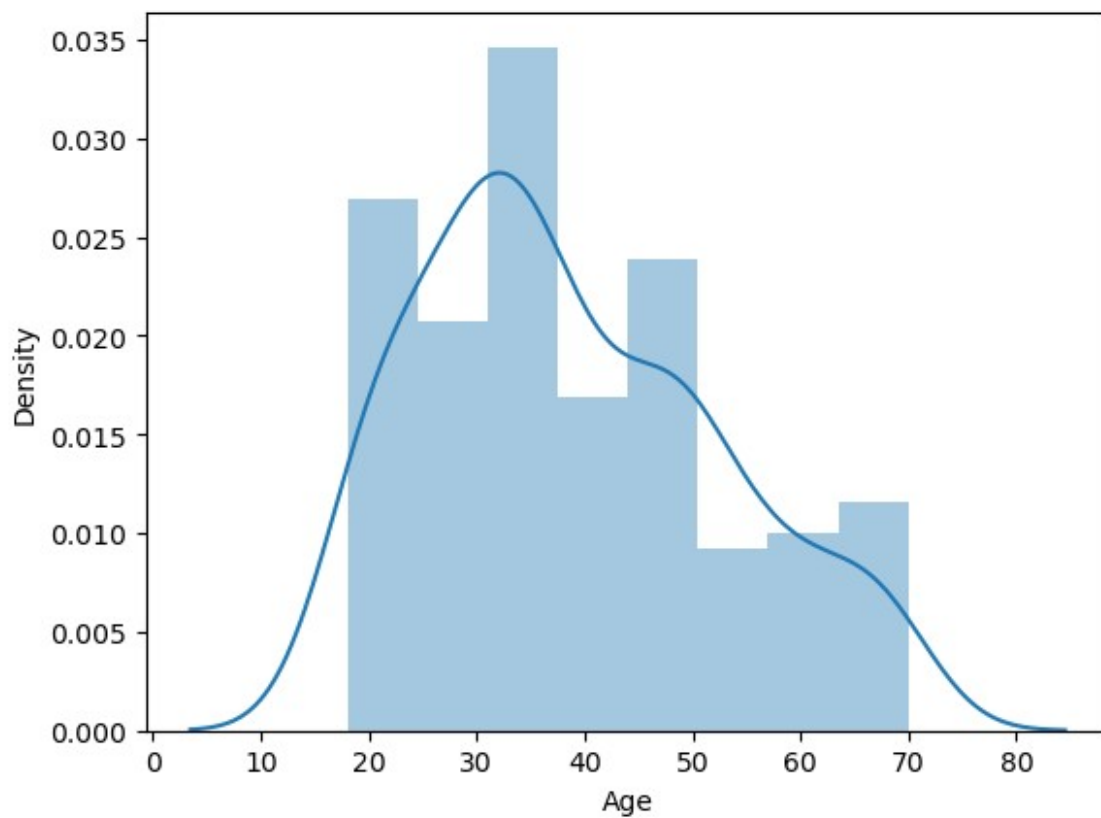
```
sns.distplot(df['Annual Income (k$)'])
```

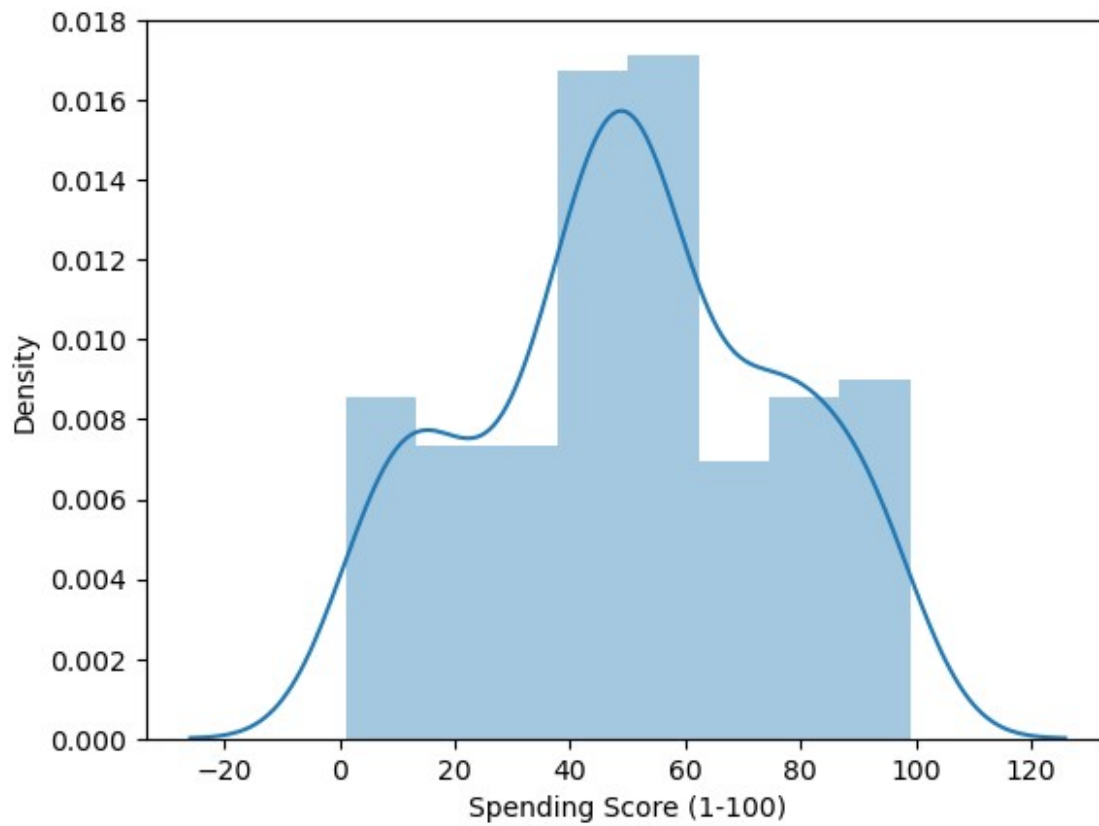
```
<Axes: xlabel='Annual Income (k$)', ylabel='Density'>
```





```
df.columns
Index(['CustomerID', 'Gender', 'Age', 'Annual Income (k$)',
      'Spending Score (1-100)'],
      dtype='object')


columns = ['Age', 'Annual Income (k$)', 'Spending Score (1-100)']
for i in columns:
    plt.figure()
    sns.distplot(df[i])
```





```
columns = ['Age', 'Annual Income (k$)', 'Spending Score (1-100)']  
for i in columns:  
    plt.figure()  
    sns.boxplot(data=df, x='Gender', y=df[i]);
```






```
df['Gender'].value_counts(normalize=True)
```

```
Gender
Female    0.56
Male      0.44
Name: proportion, dtype: float64
```

bivariate analysis

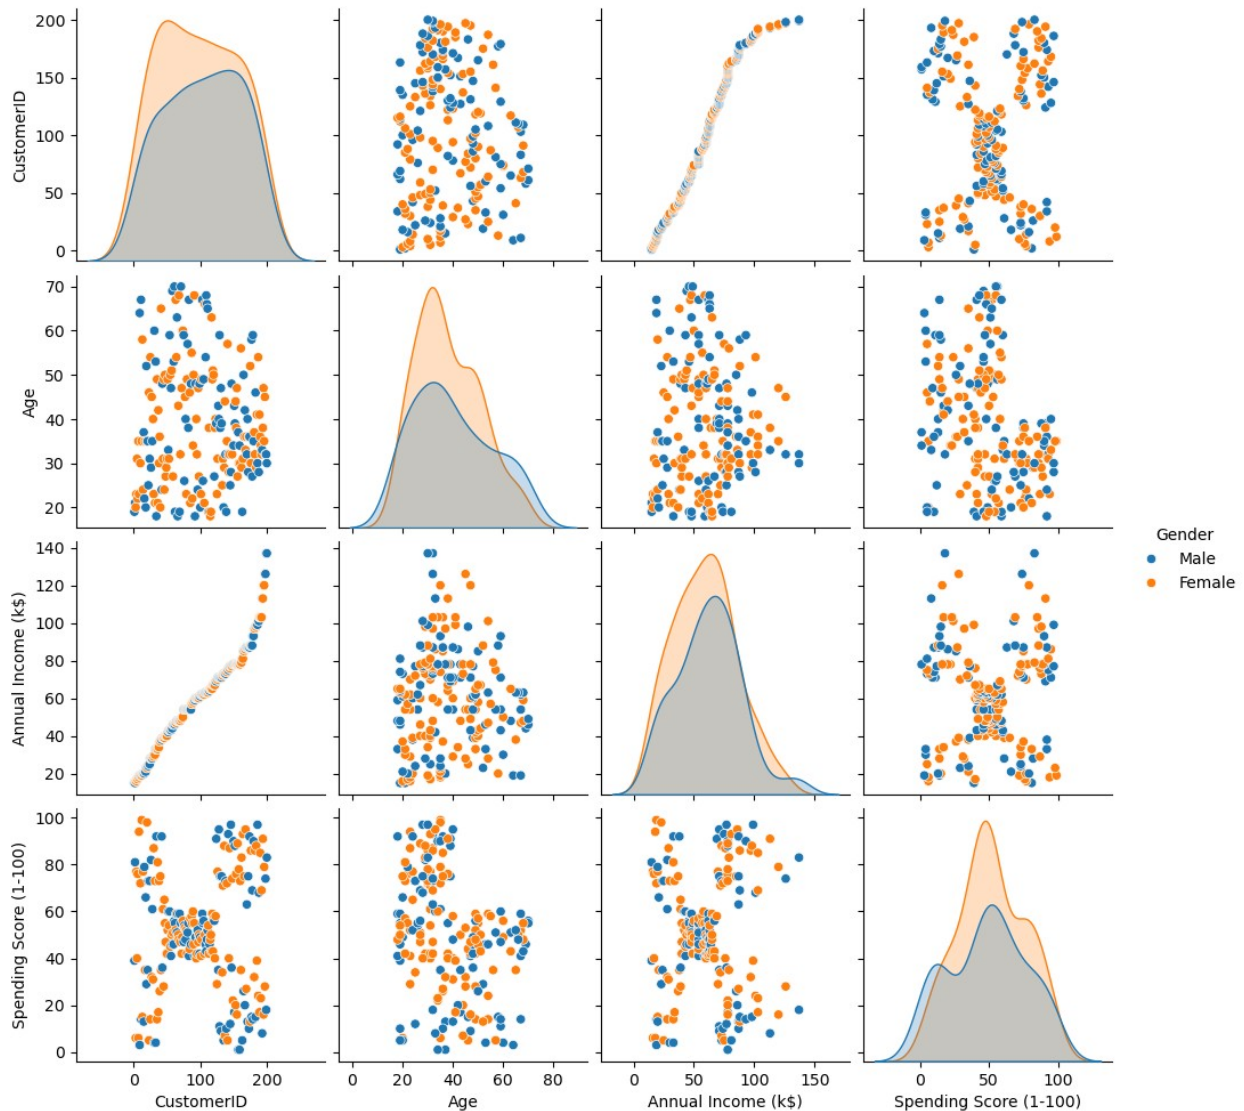
```
sns.scatterplot(data=df, x='Annual Income (k$)', y='Spending Score (1-100)')
```

```
<Axes: xlabel='Annual Income (k$)', ylabel='Spending Score (1-100)'>
```



```
#df=df.drop('CustomerID',axis=1)
sns.pairplot(df,hue='Gender')

<seaborn.axisgrid.PairGrid at 0x187b0cf23f0>
```



```
df.groupby('Gender')[['Age', 'Annual Income (k$)', 'Spending Score (1-100)']].mean()
```

	Age	Annual Income (k\$)	Spending Score (1-100)
Gender			
Female	38.098214	59.250000	51.526786
Male	39.806818	62.227273	48.511364

```
df.select_dtypes(include='number').corr()
df['Gender'] = df['Gender'].map({'Male': 1, 'Female': 0}) # Assuming
'Male' = 1 and 'Female' = 0
df.corr()
df = pd.get_dummies(df, drop_first=True) # Converts categorical
columns to dummy variables
df.corr()
```



	CustomerID	Gender	Age	Annual Income
(k\$) \				
CustomerID	1.000000	0.057400	-0.026763	
0.977548				
Gender	0.057400	1.000000	0.060867	
0.056410				
Age	-0.026763	0.060867	1.000000	
0.012398				
Annual Income (k\$)	0.977548	0.056410	-0.012398	
1.000000				
Spending Score (1-100)	0.013835	-0.058109	-0.327227	
0.009903				

	Spending Score (1-100)
CustomerID	0.013835
Gender	-0.058109
Age	-0.327227
Annual Income (k\$)	0.009903
Spending Score (1-100)	1.000000

```
sns.heatmap(df.corr(),annot=True,cmap='coolwarm')
```

```
<Axes: >
```



```
1,
    1, 1, 1, 1, 1, 1, 1, 1, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4,
4,
    4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4,
4,
    4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
2,
    2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 0, 0, 0, 0,
0,
    0, 0])
```

```
df['Income Cluster']= clustering1.labels_
df.head()
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	1	19	15	39
1	2	1	21	15	81
2	3	0	20	16	6
3	4	0	23	16	77
4	5	0	31	17	40

	Income Cluster
0	3
1	3
2	3
3	3
4	3

```
df['Income Cluster'].value_counts()
```

```
Income Cluster
```

```
4    46
1    44
5    42
3    32
2    28
0     8
```

```
Name: count, dtype: int64
```

```
from sklearn.cluster import KMeans
```

```
# Example data
```

```
X = [[1, 2], [1, 3], [2, 3], [8, 8], [8, 9], [9, 9]]
```

```
# Fit KMeans
```

```

kmeans = KMeans(n_clusters=2)
kmeans.fit(X)

# Get inertia
print(kmeans.inertia_)

2.666666666666667

intertia_scores=[]
for i in range(1,11):
    kmeans=KMeans(n_clusters=i)
    kmeans.fit(df[['Annual Income (k$)']])
    intertia_scores.append(kmeans.inertia_)


intertia_scores

[137277.280000000006,
 49761.73701298703,
 25341.285871863212,
 13278.112713472483,
 9407.908188585603,
 5742.224880382777,
 3941.4163614163635,
 2985.5603641456573,
 2459.3303030303055,
 2013.4640637140633]

plt.plot(range(1,11),intertia_scores)

[<matplotlib.lines.Line2D at 0x187b7fc4c80>]

```



```
df.columns
```

```
Index(['CustomerID', 'Gender', 'Age', 'Annual Income (k$)',  
      'Spending Score (1-100)', 'Income Cluster'],  
      dtype='object')
```

```
from sklearn.cluster import KMeans
```

```
# Assuming you are clustering based on 'Annual Income (k$)' and  
'Spending Score (1-100)'
```

```
kmeans = KMeans(n_clusters=3) # Example: 3 clusters  
df['income cluster'] = kmeans.fit_predict(df[['Annual Income (k$)',  
      'Spending Score (1-100)']])
```

```
df.groupby('income cluster')[['Age', 'Annual Income (k$)', 'Spending  
Score (1-100)']].mean()
```

	Age	Annual Income (k\$)	Spending Score (1-100)
income cluster			
0	40.325203	44.154472	49.829268
1	40.394737	87.000000	18.631579
2	32.692308	86.538462	82.128205

bivariate clustering

```
clustering2= KMeans()
clustering2.fit(df[['Annual Income (k$)', 'Spending Score (1-100)']])


df['Spending and income cluster']=clustering2.labels_
df.head()
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	1	19	15	39
1	2	1	21	15	81
2	3	0	20	16	6
3	4	0	23	16	77
4	5	0	31	17	40

	Income Cluster	income cluster	Spending and income cluster
0	3	0	3
1	3	0	4
2	3	0	6
3	3	0	4
4	3	0	3


```
intertia_scores2=[]
for i in range(1,11):
    kmeans2=KMeans(n_clusters=i)
    kmeans2.fit(df[['Annual Income (k$)', 'Spending Score (1-100)']])
    inertia_scores2.append(kmeans2.inertia_)
plt.plot(range(1,11), inertia_scores2)

[<matplotlib.lines.Line2D at 0x187b923cec0>]
```



```
centers= pd.DataFrame(clustering2.cluster_centers_)
centers.columns= ['x','y']

plt.figure(figsize=(10,8))
plt.scatter(x=centers['x'],y=centers['y'])
sns.scatterplot(data = df, x='Annual Income (k$)', y='Spending Score
(1-100)',palette='table')
plt.savefig("clustering_bivariate.png")
```



```
pd.crosstab(df['Spending and income
cluster'],df['Gender'],normalize='index')
```

Gender	0	1
Spending and income cluster		
0	0.585366	0.414634
1	0.538462	0.461538
2	0.700000	0.300000
3	0.636364	0.363636
4	0.571429	0.428571
5	0.347826	0.652174
6	0.583333	0.416667
7	0.604651	0.395349

```
df.groupby('Spending and income cluster')[['Age', 'Annual Income
(k$)', 'Spending Score (1-100)']].mean()
```

	Age	Annual Income (k\$)	\
Spending and income cluster			
0	42.024390	63.804878	



1	32.692308	86.538462
2	41.000000	109.700000
3	41.363636	28.181818
4	25.333333	25.095238
5	40.217391	79.739130
6	48.750000	24.583333
7	43.372093	47.860465

	Spending Score (1-100)
Spending and income cluster	
0	46.634146
1	82.128205
2	22.000000
3	33.272727
4	80.047619
5	13.391304
6	9.583333
7	51.976744

#multivariate clustering

```
from sklearn.preprocessing import StandardScaler
scale= StandardScaler()
df.head()
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	1	19	15	39
1	2	1	21	15	81
2	3	0	20	16	6
3	4	0	23	16	77
4	5	0	31	17	40

	Income Cluster	income cluster	Spending and income cluster
0	3	0	3
1	3	0	4
2	3	0	6
3	3	0	4
4	3	0	3

```
dff= pd.get_dummies(df,drop_first=True)
dff.head()
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	1	19	15	39
1	2	1	21	15	81
2	3	0	20	16	6
3	4	0	23	16	77
4	5	0	31	17	40

	Income Cluster	income cluster	Spending and income cluster
0	3	0	3
1	3	0	4
2	3	0	6
3	3	0	4
4	3	0	3

```
dff.columns
```

```
Index(['CustomerID', 'Gender', 'Age', 'Annual Income (k$)',
      'Spending Score (1-100)', 'Income Cluster', 'income cluster',
      'Spending and income cluster'],
      dtype='object')
```

```
print(dff.columns)
```

```
Index(['CustomerID', 'Gender', 'Age', 'Annual Income (k$)',
      'Spending Score (1-100)', 'Income Cluster', 'income cluster',
      'Spending and income cluster'],
      dtype='object')
```

```
dff['Gender_Male'] = (dff['Gender'] == 'Male').astype(int)
```

```
dff = dff[['Age', 'Annual Income (k$)', 'Spending Score (1-100)',
          'Gender_Male']]
dff.head()
```

	Age	Annual Income (k\$)	Spending Score (1-100)	Gender_Male
0	19	15	39	0
1	21	15	81	0
2	20	16	6	0
3	23	16	77	0
4	31	17	40	0

```
dff=scale.fit_transform(dff)
```

```
dff=pd.DataFrame(scale.fit_transform(dff))
dff.head()
```


	0	1	2	3
0	-1.424569	-1.738999	-0.434801	0.0
1	-1.281035	-1.738999	1.195704	0.0
2	-1.352802	-1.700830	-1.715913	0.0
3	-1.137502	-1.700830	1.040418	0.0
4	-0.563369	-1.662660	-0.395980	0.0

```

intertia_scores3=[]
for i in range(1,11):
    kmeans3=KMeans(n_clusters=i)
    kmeans3.fit(dff)
    inertia_scores3.append(kmeans3.inertia_)
plt.plot(range(1,11), inertia_scores3)

[<matplotlib.lines.Line2D at 0x187baddcda0>]

```



```

df

```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	1	19	15	
1	2	1	21	15	
2	3	0	20	16	

```

6
3          4          0    23          16
77
4          5          0    31          17
40
..          ...          ...    ...          ...
..
195        196          0    35          120
79
196        197          0    45          126
28
197        198          1    32          126
74
198        199          1    32          137
18
199        200          1    30          137
83

```

	Income Cluster	income cluster	Spending and income cluster
0	3	0	3
1	3	0	4
2	3	0	6
3	3	0	4
4	3	0	3
..	...	...	...
195	0	2	1
196	0	1	2
197	0	2	1
198	0	1	2
199	0	2	1

```
[200 rows x 8 columns]
```

```
df.to_csv("Clustering.csv")
```